
Verifying Agentic Scientific Workflows from Primitive Evidence with GRACE

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Agentic AI systems produce scientific results through long autonomous work-
2 flows, and the verification layers they carry are their own: self-assessments, plau-
3 sibility checks, and generated reports. We show these layers can carry no signal.
4 In the public artifact corpus of an autonomous detector-design agent, all 64 pre-
5 served physics-verification records return `plausible: true` with empty issue
6 lists, spanning runs whose headline results include a retyped constant reported as
7 a 204.62% improvement, a hardcoded detection efficiency of 1.0, and a place-
8 holder narrated as physics. We propose *calibrated artifact auditing*: every claim
9 is re-derived from the most primitive surviving evidence rather than the agent’s
10 intermediates or self-reports; a failure diagnosis is accepted only if it reproduces
11 the claimed number to all quoted digits and survives an empirically measured
12 false-positive rate on decoy targets; and the auditor itself is calibrated by fault in-
13 jection, shipping with measured detection profiles and a characterized blind spot
14 rather than an implicit completeness claim. A five-stage implementation audits the
15 full eighteen-run corpus (3,142 numeric claims) in under an hour on a laptop, re-
16 discovers every manually established failure mechanism from primitive evidence,
17 corrects one of them, and—via a moment-inversion technique that reconstructs
18 an analyzed file’s summary statistics from archived error formulas—exposes two
19 analyses that consumed stale data existing nowhere in the recorded provenance.
20 We state the method’s provenance preconditions and release the pipeline.

21 1 Introduction

22 Agentic systems built on large language models now execute complete scientific workflows: they
23 read papers, write and run simulation and analysis code, interpret outputs, and report conclusions
24 across dozens of autonomous steps (Hill and Ryoo, 2026a; Hellert et al., 2026; Diefenbacher et
25 al., 2025; Moreno et al., 2026; Chung et al., 2026). The question of whether their outputs can be
26 trusted is therefore no longer hypothetical, and the field’s current answer is largely *self*-verification:
27 agents carry plausibility checkers, structured self-assessments, and verification sub-agents, and their
28 demonstrations are validated by the agent’s own account of its work.

29 This paper reports a measurement that should end confidence in that arrangement, and proposes
30 what to do instead. In the public artifact release of GRACE, an autonomous agent for particle-
31 detector design and simulation (Hill and Ryoo, 2026a,b), the verification layer demonstrably fired:
32 the reasoning traces of eight runs preserve 64 structured physics-verification records. Every one of
33 them returns `plausible: true` with an empty issue list. The same corpus, audited externally,
34 contains an order-of-magnitude geometry error that the agent’s own validator rejected once and
35 then accepted on an unmodified retry, a headline “204.62% improvement” whose optimized yield
36 is a retyped baseline scaled by a retyped percentage, a detection efficiency of 100% corresponding
37 to the literal assignment `baseline_detection_efficiency = 1.0`, and a defaulted placeholder

38 whose statistics the generated report narrated as a physical effect. The run responsible for most of
39 these defects reported success with all sixteen planned steps completed and `failed_steps: 0`,
40 after a trace containing five step errors and one regeneration. Self-assessment, in this corpus, carries
41 *no signal* separating genuine results from artifacts.

42 The constructive claim is that external verification of such workflows is tractable, cheap, and—
43 critically—*calibratable*. We propose **calibrated artifact auditing**, a method defined by three com-
44 mitments:

- 45 1. **Primitive evidence only.** Every quantitative claim is traced to, and re-derived from, the
46 most primitive surviving record of the computation (here, per-event execution logs), never
47 from the agent’s intermediate results files, prose, or self-reports, which are treated as claims
48 requiring evidence rather than as evidence.
- 49 2. **Exact, null-calibrated diagnosis.** When a claim cannot be reproduced, a generating mech-
50 anism is accepted only if it reproduces the claimed number to all quoted digits *and* survives
51 an empirical false-positive measurement: the identical search rerun on jittered decoy targets
52 of equal precision must not match at chance level.
- 53 3. **Calibrate the auditor.** The audit itself is a measuring instrument. Known defects are
54 injected into clones of a real run, detection is scored per stage against rules fixed before
55 injection, and the audit ships with measured detection profiles and characterized blind spots
56 rather than an implicit claim of completeness.

57 The third commitment is, to our knowledge, new to this setting, and it is the one we most want
58 the community to adopt. A verifier evaluated only on the defects it happened to find is in the same
59 epistemic position as the agent it audits.

60 We instantiate the method as a five-stage pipeline and evaluate it on the complete eighteen-run
61 GRACE corpus. Our contributions: (i) the method and its preconditions (§3); (ii) a reference imple-
62 mentation whose expensive human steps—hypothesis formation and re-derivation—are mechanized
63 as calibrated formula search and differential re-execution (§4); (iii) a corpus-scale evaluation ac-
64 counting for 3,142 numeric claims, rediscovering all manually established mechanisms and finding
65 new defect classes, including stale-input contamination exposed by a moment-inversion technique
66 (§5); and (iv) a fault-injection calibration of the auditor itself, measuring a union-of-detectors cov-
67 erage structure with one characterized blind spot (§6).

68 2 Related work

69 Verification of AI-generated science is nascent. Execution-scored benchmarks evaluate generated
70 physics code by running it (Ji and Boriskina, 2026), capability probes measure model performance
71 on research tasks (Miao et al., 2026), and recent work argues for surfacing the assumptions inside
72 AI-generated analyses (Hammad and Nojiri, 2026). Reproducibility infrastructure in experimental
73 physics, exemplified by REANA and LHC workflow preservation (Šimko et al., 2019), answers
74 whether a preserved result can be *rerun*. The method here answers two prior questions: whether a
75 claimed result was ever computed at all, and whether it was computed from the data the workflow
76 actually produced. It requires artifacts an order of magnitude lighter than a preserved workflow,
77 and it is complementary to rerun infrastructure rather than a substitute. A companion measurement
78 addresses the orthogonal hazard of hindsight contamination in the models backing such agents.

79 3 The method and its preconditions

80 **Provenance tiers.** Every quantity is assigned a tier: LOG for values derived from per-event execu-
81 tion records (primitive evidence); JSON for numeric leaves of results files (intermediates); REPORT
82 for numeric strings in generated prose (terminal claims). Verification flows strictly upward: a RE-
83 PORT claim is supported only by tracing it to JSON and then to LOG; agreement between intermedi-
84 ates is recorded as copying, not verification.

85 **Verdicts.** Each numeric claim receives exactly one verdict: *verified* (reproduced from primitive
86 evidence by identity), *copy* (verbatim inter-intermediate duplicate), *diagnosed* (reproduced only by

87 a nontrivial mechanism passing the calibration gates below), *unexplained* (no mechanism in the
88 bounded search matches—a deliberately conservative bucket), or *unevaluable* (inputs absent from
89 the archive, with the archival gap itself recorded as a finding). Coverage is reported as fractions over
90 the full denominator, never as a list of findings.

91 **Preconditions.** The method applies to corpora satisfying a minimum provenance floor: per-event
92 (or equivalently primitive) text logs retained for every execution, and generated code retained ver-
93 batim and identifiable (e.g., by content hash). These are cheap—orders of magnitude smaller than
94 binary data products—and we argue they are the minimum any agentic scientific system should ship.
95 Two further requirements, motivated by defects found in the evaluation, extend the floor: content-
96 hashing of *input* files at read time against the workflow’s write manifest (without which stale-input
97 contamination is invisible), and an execution-status and output-attribution manifest (without which
98 content-addressed script retention cannot establish which script produced which file). Appendix D
99 tabulates the full standard against the failure class each item prevents.

100 4 Reference implementation

101 The pipeline mechanizes the two expensive steps of a manual audit—mechanism hypothesis and
102 re-derivation—and runs five stages. On the evaluation corpus it completes in under an hour on a
103 laptop with standard Python, no accelerators, and no access to the original computing environment.

104 **Stage 1: null-calibrated mechanism search.** For each REPORT claim and JSON leaf, a bounded
105 formula grammar is searched over the run’s quantity pool: identity; mean and standard deviation
106 of any numeric family, with fallback variants appending or substituting a placeholder 1.0; ratios;
107 energy-scaled ratios (the unit-loss class); percent differences; and percentage-scaled values $a(1 +$
108 $b/100)$ (the transcription-and-rescale class). The target’s own leaf is removed from the pool before
109 pair formulas run. A match must reproduce every printed digit at the claim’s precision. Three
110 guards keep the search evidential. Pair formulas are attempted only for targets with ≥ 4 significant
111 figures; matches with multiplicity > 5 are rejected as chance-level; and every surviving diagnosis is
112 assigned an empirical false-positive rate by rerunning the identical search on 100 seeded, jittered
113 decoy targets of equal precision. When mean and standard deviation of the same family variant
114 jointly match a reported value–uncertainty pair, the diagnosis is promoted to *corroborated* with a
115 joint null measured directly on jittered value–uncertainty decoy pairs, alongside the product of the
116 individual rates. Calibration cuts both ways: on the evaluation corpus it certified one diagnosis at
117 a measured joint null of zero in 100 decoy pairs while eleven single-formula matches in the same
118 run carried decoy rates of 44–99% and stayed screening verdicts. Low-precision targets match *some*
119 grammar formula most of the time (measured decoy rates versus precision in Appendix B), which is
120 precisely why uncalibrated formula-matching is inadmissible as evidence.

121 **Stage 2: dataflow analysis.** Each generated script is parsed to an abstract syntax tree and an-
122 alyzed with name-level, flow-insensitive taint propagation over literal, file-derived, and external
123 sources, with container mutations propagating taint. An output (a RESULT: print or json.dump
124 field) reachable only from literal assignments is flagged and severity-classified by label semantics
125 (physics class: yield, efficiency, resolution, quenching, significance; configuration class: counts,
126 positions, geometry). The same pass detects existence-guarded silent skips, where a file-existence
127 test dominates a continue on a path feeding an output.

128 **Stage 3: physics-invariant battery.** Checks computable from primitive evidence alone, with no
129 claim tracing: bounds on fraction-type quantities; a photostatistics floor $\sigma/\mu \gtrsim 1/\sqrt{N_{\text{PE}}}$ tested
130 against reported resolutions with N_{PE} from logs (binomial detection of a fixed photon count sits
131 slightly below the Poisson value, so the test fires only below 0.8 of the floor); a batched-run sup-
132 pression test comparing log-derived response at batched versus unbatched settings; and a linearity
133 residual scan. This is the *binding* verification layer whose absence in the agent caused the failures
134 the audit finds, demonstrated retrospectively.

135 **Stage 4: differential re-execution.** The archive excludes binary event files but retains the per-
136 event logs from which analysis inputs are reconstructible. The pipeline synthesizes the excluded
137 event files from logs, rewrites cluster-absolute paths into a sandbox, re-executes every generated

Table 1: Corpus accounting over all 18 runs, by arm. *Log-verified*: reproduced from per-event logs by identity. *Copies*: verbatim inter-file duplicates. *Diagnosed*: reproduced only by a nontrivial mechanism passing the Stage-1 gates (a screening category; only mechanisms confirmed by independent evidence are asserted as findings). *Unexplained*: no match in the bounded grammar—a conservative bucket including low-precision values and prose restatements. *HC*: physics-class script outputs reachable only from literal assignments. *Skips*: existence-guarded silent omissions. Per-run rows appear in Appendix A.

Arm (runs)	Log-ver.	Copies	Diagnosed	Unexpl.	HC	Skips
Celeritas v4 (6)	37	304	99	316	2	4
Celeritas v7 (6)	89	506	259	685	17	0
Opticks v5 (3)	21	182	28	262	24	0
Opticks v9 (3)	39	120	62	133	18	2
Total (18)	186	1112	448	1396	61	6

script under a 150s timeout, and diffs each regenerated JSON leaf against the archive. A MATCH verifies the computation end-to-end mechanically; a DIVERGE localizes a defect and names the value the correctly fed computation would have produced. One complementarity is essential: re-execution certifies that the archived script deterministically reproduces the archived output, while *derivation from data* is a separate property—a script whose outputs are literals regenerates them identically and is caught only by Stage 2. The calibration of §6 measures this complementarity directly.

Stage 5: claim accounting. Every numeric claim receives an identifier and a verdict; the corpus-level table is generated by one command. A table-row consistency check flags report rows in which multiple high-precision cells have no support anywhere in the archive.

5 Evaluation on a public agentic physics corpus

The evaluation corpus is the public GRACE benchmark release (Hill and Ryoo, 2026b): 1,248 files across 18 archived agentic runs (twelve reporting success, six ending in process-level failure) on a pinned LLM backbone, spanning two simulation frameworks (Opticks and Celeritas arms, in two framework versions each) and both paper-driven and prompt-driven tasks. Two of its properties match the method’s preconditions: the Opticks arms log one line per simulated event (the Celeritas arms log every hundredth event), and generated scripts are retained by content hash. Its gitignore excluded all binary data products, which the method tolerates by log-synthesis (Stage 4). Three results files (`projective_performance_analysis.json`, `ams_optimization_results.json`, `bgo_performance_metrics.json`) are truncated mid-object and fail to parse, so their leaves enter no column of the accounting; the truncation pattern, a key with no value at the end of the file, is the signature of a serializer raising partway through a dump. A prior manual audit of the two paper-driven runs, requiring several person-days, established a set of failure mechanisms by hand; it serves as the ground-truth reference for whether the mechanized audit recovers known defects without hypotheses supplied.

5.1 Claim accounting

The pipeline accounts for 3,142 numeric claims (Table 1); in the mechanized accounting, targets whose inputs are absent from the archive are not tracked as a separate column and fall within the conservative unexplained bucket. Three observations matter for verification practice. First, only 186 claims (6%) are verifiable against primitive evidence by identity, while 1,112 (35%) are copies between intermediates: the corpus *looks* heavily cross-confirmed while most agreement is duplication, which is exactly the structure that defeats consistency-based checking (§6). Second, 61 physics-class outputs across ten runs are reachable only from literal constants, including a script that computes per-energy quenching factors in a loop, discards them, and writes a hand-transcribed literal in their place under the comment `# Use values from previous step analysis`. We name this *transcription hardcoding*: the agent’s stdout-mediated memory between steps makes transcription the path of least resistance, and nothing in a results file distinguishes a transcribed constant from a com-




Figure 1: Stale-input contamination in both paper-driven runs: mean photoelectrons per event vs. beam energy. Circles: log-verified points, with a linear response fit (dashed). Squares: the genuine 5 MeV data computed from the archived logs, suppressed to 0.47 of the linear response by the batching artifact of §5.2. Open stars: the 5 MeV entries the archives actually record, plotted at the mean implied by moment inversion. In both runs the archived entry sits on the linear response of the genuine points and carries exactly the intended 5000 events, which every logged sweep missed—so the analyzed files exist nowhere in the recorded provenance.

putation. Third, the unexplained bucket (44%) is reported as such: a bounded, calibrated search that cannot certify a derivation says so, rather than absorbing ambiguity into apparent coverage.

5.2 Rediscovery of known mechanisms

Without hypotheses supplied, the pipeline rediscovered from primitive evidence every quantitative mechanism the manual audit had established, and replaced the manual reading of one of them with an exact-digit derivation (complete arithmetic in Appendix C): the headline “204.62% improvement” as a retyped constant, with the optimized yield 4491.29 reproduced to all stored digits as the cross-energy baseline mean 1474.39 scaled by that constant (the log-reconstructed true result is a uniform $\sim 28.6\%$ improvement, physically sensible and sublinear in the 33.3% coverage increase); the reported quenching mean and standard deviation (0.631, 0.264) as the statistics of two log-derived values plus a placeholder 1.0, corroborated at a measured joint decoy null of zero in 100 decoy pairs; a systematic yield suppression at exactly the energy where execution was batched, appearing in three run-particle combinations across two unrelated geometries at factors 0.47, 0.47, and 0.22; and the exact existence-check responsible for a silent input omission. The generated report had narrated the placeholder’s statistics as a convergence to unity at 5 MeV, offered as either an artifact or a physical effect; the agent’s self-assessment recorded sixteen completed steps and zero failed steps. The geometry default of the introduction lives in the extraction layer, upstream of any number the five stages consume, and remains a source-inspection finding.

5.3 Case study: stale-input contamination by moment inversion

One archived value had resisted the manual audit: a 5 MeV light yield matching no quantity derivable from any log. The pipeline closes it with a technique we believe is new in this setting. The four archived leaves for the entry overdetermine the moments of whatever file the analysis read, because the script’s error formulas are archived verbatim: $\sigma_R = R/\sqrt{2n}$ applied to the archived ratio yields $n = 5000$ events exactly; yield and resolution then fix $\mu = 7430.8254$ and $\sigma = 119.0499$; and the fourth leaf follows with no remaining freedom, matching the archive to all stored digits. The analyzed file’s complete summary statistics are therefore known—and they match nothing in the recorded provenance. No log of any geometry, particle, or energy has a mean within 5% of the reconstruction. A one-line bound already settles the baseline-electron stream, since the largest per-event value in any 5 MeV baseline-electron log is 3,969, below the reconstructed mean of 7,430.8, so no subset, prefix, or suffix of that stream can average it. Exhaustive enumeration over every batch subset and every event-level prefix and suffix of both particle streams (the complete hypothesis space of partial merges, truncations, and a filename collision found by Stage 4) yields zero matches at four-leaf precision. The value was honestly computed, by the archived script, from a

Table 2: Fault-injection calibration, instances detected out of three per class, with every stage run on every instance (75 measured cells). Each class is caught by at least one stage, each stage has a distinct catch profile, and no stage subsumes another.

Injected defect	X-check	Dataflow	Coverage	Mech. search	Re-exec.
D1 JSON digit flip	1/3	0/3	0/3	0/3	3/3
D2 hardcoded physics output	0/3	3/3	0/3	0/3	0/3
D3 log deletion	0/3	0/3	3/3	1/3	0/3
D4 fabricated report number	3/3	0/3	0/3	3/3	0/3
D5 placeholder-mean statistic	0/3	0/3	0/3	3/3	0/3

stale file at the expected path—data outside the recorded workflow. The same $n = 5000$ fingerprint appears independently in the second paper-driven run, where the archive writes `num_events`: 5000 outright against 4,944 logged events, revising a claim the manual audit had *accepted* as verified (Fig. 1). Stale-input contamination is invisible to any report- or results-level check, because the numbers are computed correctly from the wrong data; it yields only to input-content attestation at write/read time, or to moment inversion, which converts a handful of reported decimals into a forensic fingerprint of the input file.

5.4 Differential re-execution

Of 46 generated analysis scripts in the Opticks arms, 25 re-executed in the sandbox (the remainder read unreconstructible inputs or are archive-defective; two archived scripts respectively terminate mid-string and call `json.dump` with no file object, so one can never have produced the results file bearing its name, and two scripts in the Celeritas-v7 Higgs run assign `df = null`, a `NameError` at first use). Twenty-one regenerated their results files in exact leaf-for-leaf agreement. Each of the four divergences is individually informative: two localize the stale inputs of §5.3, one, the script that computes the 204.62% figure of §5.2, collapses to zeros when its existence-guarded inputs are absent and demonstrates input-path fragility, and one diverges purely through seedless random number generation.

5.5 The self-assessment layer, quantified

Across the corpus, twelve of eighteen runs record `success: true`; all six failures are process-level (three iteration limits, three incomplete planned steps), none semantic. All 64 preserved physics-verification records return `plausible: true` with empty issue lists—unanimity spanning runs containing every defect in this section. The corpus also shows what a binding check achieves on its own. The geometry validator of the DarkSide run rejected the first geometry as orders of magnitude larger than the DarkSide-50 specification, forced a regeneration, received a retry issued with an empty modification set, and returned `VALID` on a vessel of 3.6 m by 3.6 m against a real instrument near 0.36 m, with the extraction still carrying no TPC dimensions. A verification layer that binds once and then accepts the same defect on retry, or that never writes outcomes into results files, adds little to auditability; one that passes everything is worse than none, because it lends artifacts an imprimatur of scrutiny. The invariant battery of Stage 3 shows constructively that a binding check was available to the agent at run time: every input to the batch-suppression test that flags the 0.47 artifact was present in the agent’s own logs.

6 Calibrating the auditor by fault injection

An audit that reports findings without a detection-rate estimate cannot answer the question every reviewer should ask: how does the auditor know the search was complete? We mutated clones of one run with five defect classes, three independent instances of each, with detection rules fixed before injection, scoring each stage per instance; unmutated control clones establish each detector’s false-alarm floor.

The blind spot is the calibration’s most useful output. Single-leaf corruption (D1) evades every detector that asks whether a reported number matches *some* archived value whenever the corrupted leaf

has a duplicate, two of the three instances here, because value redundancy across results files—the 35% copy rate of Table 1—means the corrupted leaf coexists with uncorrupted duplicates; the cross-check caught only the instance whose leaf had no duplicate, and regeneration from primitive inputs exposed all three. The complementarity runs the other way as well: the transcription-hardcoded script re-executes to an *exact match*, because literals regenerate identically, and is caught only by dataflow. Two off-diagonal catches are informative as well, since a fabricated report number also fails the plain cross-check, and deleting a log turns one formerly log-verified claim unexplained. Coverage is therefore the union of five differently shaped detectors, none subsuming another, and the method ships with that structure *measured* rather than asserted. We propose this as a norm: a verifier for agentic science should publish its fault-injection profile the way a physical instrument publishes its response function.

7 Limitations and outlook

The method is evaluated on one system’s corpus, and its generality is a proposed, testable claim: the mechanisms it detects live overwhelmingly in generated pipeline code (unit loss, existence-guarded skips, hardcoded constants, stale paths) rather than in model-specific behavior, which predicts transfer across agents and backbones, and the falsifiable form of the prediction is that fault injection on any agentic corpus meeting the preconditions will reveal the same union-of-detectors structure. The taint analysis is name-level and flow-insensitive—spot-verified on every reported finding class, but formally unsound. The fault-injection set covers five defect classes and cannot bound classes outside it. The Celeritas arms log only one event in a hundred, so re-execution there certifies only non-event-derived outputs. Prose–archive drift detection is string-based and misses unit-converted restatements, making the claim denominator conservative. The moment-inversion results condition on archived scripts matching executed ones, confirmed here by byte-identical trace output.

Two directions follow. First, the preconditions are an actionable standard: agentic systems should ship per-event logs, content-hashed code, input-content attestation, and output-attribution manifests, at negligible cost relative to data volume. Second, the invariant battery demonstrates that binding, physics-aware checks were computable from information the agent already had; moving such checks inside agents, with the authority to halt, converts this external audit into a run-time discipline. The pipeline, the frozen stage outputs reproducing every number here, and the fault-injection harness are released for blind reruns by uninvolved parties.¹

References

- J. Hill and H. J. Ryoo. GRACE, an agentic AI for particle physics experiment design and simulation. arXiv:2602.15039, 2026.
- J. Hill and H. J. Ryoo. GRACE whitepaper data release. https://github.com/just5034/GRACE_whitepaper_data, 2026.
- T. Hellert, D. Bertwistle, S. C. Leemann, A. Sulc, and M. Venturini. Agentic artificial intelligence for multistage physics experiments at a large-scale user facility particle accelerator. Phys. Rev. Research 8, L012017, 2026.
- S. Diefenbacher, A. Hallin, G. Kasieczka, M. Krämer, A. Lauscher, and T. Lukas. Agents of discovery. arXiv:2509.08535, 2025.
- E. A. Moreno, S. Bright-Thonney, A. Novak, D. Garcia, P. Harris, et al. AI agents can already autonomously perform experimental high energy physics. arXiv:2603.20179, 2026.
- W. Chung, Q. Liu, L. Wu, and J. Gonski. Agentic-AI detector co-design and optimization in vertically-integrated differentiable full simulations. arXiv:2604.21804, 2026.
- X. Ji and S. V. Boriskina. PHITSBench, an execution-scored benchmark for AI-assisted PHITS radiation-transport input generation using natural language. arXiv:2607.09789, 2026.
- T. Miao, W. Jin, M. Zhang, J. Tan, Y. Hu, et al. PRL-Bench, a comprehensive benchmark evaluating LLMs’ capabilities in frontier physics research. arXiv:2604.15411, 2026.

¹Anonymized review artifact: https://anonymous.4open.science/r/grace_audit-EC02. It contains the five-stage pipeline (artifact `grace_audit_adfa1a25`, named by its pipeline source hash), the frozen per-stage outputs regenerated by that version, and the figure-generation scripts; the audited corpus itself is fetched separately per its README.

- 293 A. Hammad and M. Nojiri. Articulating assumptions in AI-generated scientific analyses through task decomposition.
 294 arXiv:2607.05762, 2026.
- 295 T. Šimko, L. Heinrich, H. Hirvonsalo, D. Kousidis, and D. Rodríguez. REANA, a system for reusable research data analyses.
 296 EPJ Web Conf. 214, 06034, 2019.

297 A Per-run claim accounting

Table 3: Per-run accounting underlying Table 1.

Run	Log-ver.	Copies	Diag.	Unexpl.	HC	Skips
celeritas-v4 calorimeter	7	36	33	38	0	0
celeritas-v4 emcal	14	102	28	64	0	0
celeritas-v4 higgs	0	14	0	61	0	1
celeritas-v4 muon	3	34	23	26	0	2
celeritas-v4 tof	8	6	2	89	0	1
celeritas-v4 tracking	5	112	13	38	2	0
celeritas-v7 calorimeter	10	55	20	98	4	0
celeritas-v7 emcal	3	175	66	129	3	0
celeritas-v7 higgs	0	6	6	102	0	0
celeritas-v7 muon	16	160	56	64	3	0
celeritas-v7 tof	60	77	89	259	7	0
celeritas-v7 tracking	0	33	22	33	0	0
opticks-v5 darkside	5	76	12	199	11	0
opticks-v5 detector	12	39	4	35	5	0
opticks-v5 protodune	4	67	12	28	8	0
opticks-v9 darkside	25	68	51	50	14	2
opticks-v9 detector	8	12	8	43	0	0
opticks-v9 protodune	6	40	3	40	4	0
Total	186	1112	448	1396	61	6

298 B Null calibration of the mechanism search




Figure 2: False-diagnosis rate of the pair-formula search on decoy targets as a function of target precision, measured on one run’s quantity pool with 120 decoys per bin. Low-precision targets match some formula in the grammar most of the time, motivating the gate at four significant figures; the residual rate near one half at the gate boundary is why every surviving diagnosis additionally carries its own decoy-measured null rate and multiplicity cap.

299 C Detailed mechanism arithmetic

Table 4: Headline claims of the paper-driven runs tracing exactly to simulation output.

Claim	Reported	Log-derived	Primary evidence
Run A light yield at 1 MeV	1479.23 PE/MeV	1479.2	7.40M hits / 5k events
Run A resolution at 1 MeV	0.0258	0.0258	per-event hit distribution
Quenching at 0.1, 1 MeV	0.396, 0.498	0.3964, 0.4978	proton/electron yield ratios
Run B yield at 1, 2 MeV	651.4, 653.3 PE/MeV	651.4, 653.3	5,000 events per energy

300 **Verified-claim arithmetic.** Run A (DarkSide-style geometry) at 1 MeV: 7,396,169 hits over 5,000
301 events = 1479.2 PE/event = PE/MeV at 1 MeV. Quenching at 0.1 MeV: $577.8/1457.8 = 0.3964$;
302 at 1 MeV: $736.4/1479.2 = 0.4978$. Run B (ProtoDUNE-style geometry): 651.4 and 653.3 PE/MeV
303 from 5,000 logged events at each energy.

304 **Transcription-scaled optimized figure.** The archived comparison script assigns
305 `light_yield_improvement = 204.62` and `baseline_light_yield = 1474.39` as liter-
306 als and writes the optimized yield as $1474.39 \times (1 + 204.62/100) = 4491.286818$, which
307 reproduces the archived leaf 4491.2868180000005 to every stored digit with a unique grammar
308 match at a decoy null of 0%. The percentage originates one script upstream as total hit-record rows
309 divided by total deposited energy, summed over whichever hit files passed an existence guard, from
310 inputs absent from the archive. The manual audit had read 4491.29 as a per-event quantity with a
311 lost energy division, on the strength of $22,200,266/4491.29 = 4,942.96 \approx 4,948$ events, a match
312 at 0.1% that fails the all-digits gate. Baseline figure: $\text{mean}(1457.76, 1479.23, 1486.17) = 1474.39$
313 to all digits, a cross-energy mean whose third entry is the stale-input value of §5.3. The correct
314 optimized 5 MeV yield from the logs is $22,200,266/4,948/5 = 897.3$ PE/MeV. Reconstructed from
315 the logs on equal terms (electron vs. electron at matched energies), the true improvement is 28.7,
316 28.6, 28.4% at 0.1, 1, 5 MeV.

317 **Quenching placeholder.** $\text{mean}(0.3964, 0.4978, 1.0) = 0.6314$; population SD = 0.2639; both
318 match the report to all quoted digits. The log-derived 5 MeV quenching is $160.7/699.0 = 0.230$, so
319 the third value is a defaulted placeholder, not physics.

320 **Batching suppression.** 5-to-1 MeV baseline yield ratios: Run A $699.0/1479.2 = 0.47$; Run B
321 $307.3/651.4 = 0.47$; Run A protons = 0.22. The suppression coincides with photon-budget batch-
322 ing applied only at 5 MeV.

323 D Minimum viable provenance checklist

Table 5: Minimum viable provenance for autonomous simulation agents: the preconditions of §3 in checklist form, each item tied to the failure class it prevents or makes mechanically detectable.

Requirement	Failure prevented
Per-event text logs with defined fields, retained for every execution	enables all primitive-evidence verification
Content-hash every <i>input</i> file at read time; verify against the workflow’s write manifest	stale-input contamination (§5.3); sweep collisions
Execution-status and output-attribution manifest alongside content-addressed scripts	orphaned results files
Unique, parameter-derived output paths per geometry, particle, energy, batch	silent overwrites
Seed and record all stochastic sampling	nondeterministic reruns
Binding invariant checks (bounds, Poisson floor, batch-ratio, linearity) that <i>fail the step</i> rather than annotate it	the batching suppression artifact
Verification outcomes written into results files; report numbers quoted by reference, never retyped	transcription hardcoding; report–archive drift

324 **E Verdict definitions and search-gate parameters**

325 Pair formulas in the Stage-1 grammar are attempted only for targets with at least four significant
326 figures; matches with multiplicity above five are rejected as chance-level; decoy calibration uses
327 100 jittered targets of equal precision per diagnosis, seeded from the target text, drawn from the
328 same run's quantity pool. The unexplained verdict absorbs, by construction: low-precision val-
329 ues below the pair-formula gate, prose restatements of archived values at reduced precision, and
330 multiplicity-rejected ambiguous matches. Diagnosed is a screening verdict; promotion to an as-
331 serted finding requires independent confirmation by source inspection, re-execution, or corroborated
332 value-uncertainty statistics.